

Statistical Modeling 2

Exercise 2

February 11, 2017

1 A simple Gaussian location model

A

The joint prior over the mean parameter θ and precision parameter ω is:

$$p(\theta, \omega) \propto \omega^{(d+1)/2-1} \exp \left\{ -\omega \frac{\kappa(\theta - \mu)^2}{2} \right\} \exp \left\{ -\omega \frac{\eta}{2} \right\}$$

To get the marginal prior, we integrate out the parameter ω :

$$\begin{aligned} p(\theta) &\propto \int_0^\infty \omega^{(d+1)/2-1} \exp \left\{ -\omega \frac{\kappa(\theta - \mu)^2 + \eta}{2} \right\} \\ &\propto \left(\frac{\kappa(\theta - \mu)^2 + \eta}{2} \right)^{-(d+1)/2} \\ &= \left(\frac{\eta}{2} + \frac{\kappa(\theta - \mu)^2}{2} \right)^{-(d+1)/2} \\ &= \left(1 + \frac{\kappa(\theta - \mu)^2}{\eta} \right)^{-(d+1)/2} \left(\frac{\eta}{2} \right)^{-(d+1)/2} \\ &\propto \left(1 + \frac{\kappa(\theta - \mu)^2}{\eta} \right)^{-(d+1)/2} \\ &= \left(1 + \frac{1}{d} \frac{\kappa(\theta - \mu)^2}{\eta} \right)^{-(d+1)/2} \end{aligned}$$

Let $\nu = d, m = \mu$ and $s = \sqrt{\eta/\kappa}$, we have a Student t distribution with ν degrees of freedom and scale s :

$$p(\theta) \propto \left(1 + \frac{1}{\nu} \frac{(\theta - m)^2}{s^2} \right)^{-(\nu+1)/2}$$

B

The sampling model is:

$$(y_i \mid \theta, \omega) \sim N(\theta, 1/\omega)$$

where y_1, \dots, y_n are the datapoints, θ is the mean and ω is the precision. We have that the likelihood for all the datapoints can be written as:

$$\begin{aligned} p(\mathbf{y} \mid \theta, \omega) &\propto \prod_{i=1}^n \omega^{1/2} \exp \left\{ -\frac{1}{2} \omega (y_i - \theta)^2 \right\} \\ &= \omega^{n/2} \exp \left\{ -\frac{1}{2} \omega \sum_{i=1}^n (y_i - \theta)^2 \right\} \\ &= \omega^{n/2} \exp \left\{ -\frac{1}{2} \omega \left(\sum_{i=1}^n y_i^2 + \sum_{i=1}^n \theta^2 - 2 \sum_{i=1}^n y_i \theta \right) \right\} \\ &= \omega^{n/2} \exp \left\{ -\frac{1}{2} \omega \left(\sum_{i=1}^n y_i^2 + n\theta^2 - 2n\bar{y}\theta + n\bar{y}^2 - n\bar{y}^2 \right) \right\} \end{aligned}$$

where $\bar{y} = (\sum_{i=1}^n y_i) / n$. Let $S_y = \sum_{i=1}^n (y_i - \bar{y})^2$, we have:

$$\begin{aligned} S_y &= \sum_{i=1}^n y_i^2 + n\bar{y}^2 - 2 \sum_{i=1}^n y_i \bar{y} \\ &= \sum_{i=1}^n y_i^2 - n\bar{y}^2 \end{aligned}$$

Therefore, the likelihood is:

$$\begin{aligned} p(\mathbf{y} \mid \theta, \omega) &= \omega^{n/2} \exp \left\{ -\frac{1}{2} \omega [S_y + n(\theta^2 - 2\bar{y}\theta + \bar{y}^2)] \right\} \\ &= \omega^{n/2} \exp \left\{ -\frac{1}{2} \omega [S_y + n(\bar{y} - \theta)^2] \right\} \end{aligned}$$

The posterior is proportional to the product of the likelihood and the prior:

$$\begin{aligned} p(\theta, \omega \mid \mathbf{y}) &\propto \omega^{(d+1)/2-1} \exp \left\{ -\omega \frac{\kappa(\theta - \mu)^2}{2} \right\} \exp \left\{ -\omega \frac{\eta}{2} \right\} \\ &\quad \omega^{n/2} \exp \left\{ -\frac{1}{2} \omega [S_y + n(\bar{y} - \theta)^2] \right\} \\ &= \omega^{(d+n+1)/2-1} \exp \left\{ -\omega \frac{\kappa(\theta - \mu)^2 + n(\theta - \bar{y})^2}{2} \right\} \exp \left\{ -\omega \frac{\eta + S_y}{2} \right\} \end{aligned}$$

We also have:

$$\begin{aligned}
\kappa(\theta - \mu)^2 + n(\theta - \bar{y})^2 &= \kappa\theta^2 + \kappa\mu^2 - 2\kappa\theta\mu + n\theta^2 + n\bar{y}^2 - 2n\theta\bar{y} \\
&= (\kappa + n)\theta^2 - 2\theta(\kappa\mu + n\bar{y}) + (\kappa\mu^2 + n\bar{y}^2) \\
&= (\kappa + n) \left(\theta^2 - 2\theta \frac{\kappa\mu + n\bar{y}}{\kappa + n} + \frac{(\kappa\mu + n\bar{y})^2}{(\kappa + n)^2} \right) - \frac{(\kappa\mu + n\bar{y})^2}{\kappa + n} + (\kappa\mu^2 + n\bar{y}^2) \\
&= (\kappa + n) \left(\theta^2 - \frac{\kappa\mu + n\bar{y}}{\kappa + n} \right)^2 - \frac{(\kappa\mu + n\bar{y})^2}{\kappa + n} + (\kappa\mu^2 + n\bar{y}^2)
\end{aligned}$$

Therefore, the posterior is:

$$p(\theta, \omega \mid \mathbf{y}) \propto \omega^{(d^*+1)/2-1} \exp \left\{ -\omega \frac{\kappa^*(\theta - \mu^*)^2}{2} \right\} \exp \left\{ -\omega \frac{\eta^*}{2} \right\}$$

where:

$$\begin{aligned}
d^* &= d + n \\
\kappa^* &= \kappa + n \\
\mu^* &= \frac{\kappa\mu + n\bar{y}}{\kappa + n}
\end{aligned}$$

and

$$\begin{aligned}
\eta^* &= \eta + S_y - \frac{(\kappa\mu + n\bar{y})^2}{\kappa + n} + (\kappa\mu^2 + n\bar{y}^2) \\
&= \eta + S_y + \frac{(\kappa + n)(\kappa\mu^2 + n\bar{y}^2) - \kappa^2\mu^2 - n^2\bar{y}^2 + 2\kappa n\mu\bar{y}}{\kappa + n} \\
&= \eta + S_y + \frac{\kappa n\mu^2 + \kappa n\bar{y}^2 + 2\kappa n\mu\bar{y}}{\kappa + n} \\
&= \eta + S_y + \frac{\kappa n(\mu + \bar{y})^2}{\kappa + n}
\end{aligned}$$

C

The conditional distribution is:

$$p(\theta \mid \mathbf{y}, \omega) \propto \exp \left\{ -\omega \frac{\kappa^*(\theta - \mu^*)^2}{2} \right\}$$

We see that this is a Normal distribution with mean μ^* and variance $1/(\omega\kappa^*)$.

D

The marginal posterior of ω is:

$$\begin{aligned}
p(\omega \mid \mathbf{y}) &= \int_{-\infty}^{\infty} p(\omega, \theta \mid \mathbf{y}) d\theta \\
&\propto \omega^{(d^*+1)/2-1} \exp\left\{-\omega \frac{\eta^*}{2}\right\} \int_{-\infty}^{\infty} \exp\left\{-\omega \frac{\kappa^*(\theta - \mu^*)^2}{2}\right\} d\theta \\
&\propto \omega^{(d^*+1)/2-1} \exp\left\{-\omega \frac{\eta^*}{2}\right\} \quad (\text{Gaussian integral})
\end{aligned}$$

We see that this marginal is a Gamma distribution with parameter $(d^*/2, \eta^*/2)$.

E

The marginal posterior of θ is:

$$\begin{aligned}
p(\theta \mid \mathbf{y}) &= \int_0^{\infty} p(\theta, \omega \mid \mathbf{y}) d\omega \\
&= \int_0^{\infty} \omega^{(d^*+1)/2-1} \exp\left\{-\omega \frac{\kappa^*(\theta - \mu^*)^2 + \eta^*}{2}\right\} d\omega
\end{aligned}$$

This is the same integral in part A. By the results in A, we can see that this marginal is a Student t distribution with parameters $\nu = d^*, m = \mu^*$ and $s = \sqrt{\eta^*/\kappa^*}$.

F

FALSE. As κ approaches 0, the Normal prior on θ approaches a point distribution but the density at that point is infinite. As d and η approach 0, the Gamma prior on ω also approach a point distribution with infinite density.

G

By the results in D and E, we see that when the prior parameters approach 0, the posterior parameters are not 0 then $p(\theta \mid \mathbf{y})$ and $p(\omega \mid \mathbf{y})$ are valid distribution.

H

The classical frequentest confidence interval for θ is:

$$\bar{y} \pm t^* \frac{\sqrt{S_y}}{\sqrt{n(n-1)}}$$

As the prior parameters κ, d, η approach 0, we have the Bayesian credible interval for θ is:

$$m \pm t^* s$$

from the results in B and E, we have

$$m = \mu^* = \bar{y}$$

and

$$s = \sqrt{\eta^*/\kappa} = \sqrt{\frac{S_y + \bar{y}/n}{n}}$$

We see that this is different from the classical confidence interval.

2 The conjugate Gaussian linear model

A

The conditional posterior is proportional to the conditional prior times the likelihood:

$$\begin{aligned} p(\beta \mid \mathbf{y}, \omega) &\propto p(\beta \mid \omega) p(\mathbf{y} \mid \beta, \omega) \\ &= N(\beta; m, (\omega K)^{-1}) N(\mathbf{y}; X\beta, (\omega \Lambda)^{-1}) \\ &\propto \omega^{1/2} K^{1/2} \exp \left\{ \frac{1}{2} (\beta - m)^T \omega K (\beta - m) \right\} \\ &\quad \omega^{1/2} \Lambda^{1/2} \exp \left\{ \frac{1}{2} (\mathbf{y} - X\beta)^T \omega \Lambda (\mathbf{y} - X\beta) \right\} \end{aligned}$$